

# CSCI 3485 Lab 2

Rafael Almeida, Ryan Novitski

September 19, 2024

## 1 Introduction

Neural networks are powerful machine learning tools that can be used to solve complex data classification problems using flexible and adaptive architectures. However, their performance heavily relies on a plethora of complex parameters that directly impact its effectiveness, such as the number of layers, the number of units per layer, the learning rate, and the optimization algorithm. Understanding how these parameters effect the overall design is crucial for its success. Our goal for this lab is to explore these parameters and provide a comprehensive analysis on how these parameters affect the model's performance for the purpose of better understanding when choosing these values in future models.

## 2 Methodology

The first experiment is on 2D floating point values. Two algorithms were used to classify the points into centered blobs, and anisotropically distributed points. The experimental variables for the neural networks were the number of layers and units per layer. Additionally, the networks were trained on data with differing number of classes. This experiment is relevant because it will tell us the affects of increasing layers and units on the classification accuracy, and how this affects the classification score with differing data distributions and number of classes. The hypothesis is that increasing the complexity of the network architecture should improve its classification accuracy. Additionally, the number of classes should decrease the accuracy of the networks, but this may vary in degrees given the network architecture. Moreover, it is hard to predict the effects different data distributions will have on the classification accuracy of the networks.

The second experiment uses the same type of data and clustering algorithm. It tests the effect of learning rates and learning rate strategies on the effectiveness of the classification model. Using the same datasets as the tests above, the neural network is trained using varying learning rates from 0.001 to 0.991, with each of the three different learning strategies available. The lower the learning rate is, the less it converges each layer. Because of this, various strategies regarding the learning rate were formed to attempt to balance under/overfitting data by jumping too little or too far. The hypothesis is that the more dynamic

the learning rate strategy is, the better it will perform. Additionally, lower learning rates will generally be more accurate than higher rates.

The third experiment tests the effect different solvers have on the effectiveness of the classification model. Using the same datasets and learning rates as the tests above, the neural network is tested using both the SGD and ADAM solvers. While SGD is a simple and effective approach, ADAM is more dynamic and better able to prevent under/overfitting of the data. The hypothesis is that the ADAM solver will outperform the SGD solver in all categories.

## 3 Results


### 3.1 Layers and Units

The parameters for this experiment are:


- Number of data points: 500
- Number of features: 2
- Random state: 1
- Distribution algorithms: blob and anisotropic
- Range of number of classes 2-6
- Range of number of layers: 1-40
- Range of number of units per layer: 1-40

The number of data points is set to 500 because larger values would take much longer to train and would not change the general direction of the results besides the classification accuracy magnitude, as seen in Figure 1. The random state was set to 1 because experimentation demonstrated that it would create more complex data to classify linearly. Therefore, a value of 1 might result in more interesting data sets to attempt to classify. A range of 40 for the number of units per layer was chosen because, around this value, the neural networks would begin to perform horribly. Any larger values would yield only a few interesting results, as seen in Figure 2.

### 3.1.1 The Effects of Data Size



(a) Data size of 500



(b) Data size of 5000

Figure 1: Notice the data size does not affect the trend

### 3.1.2 Going Above 40 Units

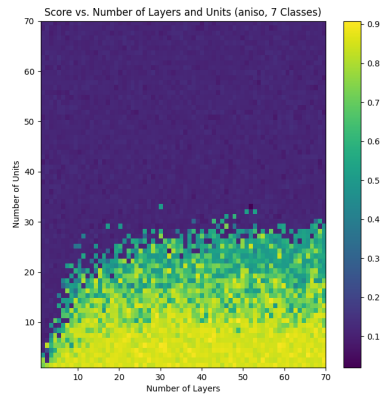








Figure 2: Notice the classification accuracy is low for any value greater than 40 units per layer

### 3.1.3 Score vs. Number of Layers and Units For Varying Data Distributions and Number of Classes







### 3.2 Solver Choice and Learning Rate

The parameters for this experiment are:

- Number of data points: 1000
- Number of features: 2
- Random state: 1
- Distribution algorithms: blob and anisotropic
- Range of number of classes: 2-3
- Number of layers: 10
- Number of units per layer: 5
- Range of initial learning rates: 0.001-0.991 (0.01 step size)
- Learning rate types: Constant, invscaling, adaptive
- Solvers used: SGD, ADAM

The number of datapoints is arbitrarily fixed at 1000. To maintain consistency with previous tests, the number of features, random state, and plot distributions remain the same. The range of number of classes was lowered to 2-3 to provide better visualizations for the data using different solvers. Based on the results of the previous tests, the highest efficiency was found with 10+ layers and 1-5 units per layer; to account for this, this experiment sets the layers and units at 10 and 5 respectively. All initial learning rates between .001 and .991 are tested at .01 increments using both solvers and all three variations of the learning rate.

### 3.2.1 Effects of Solver Choice: SGD

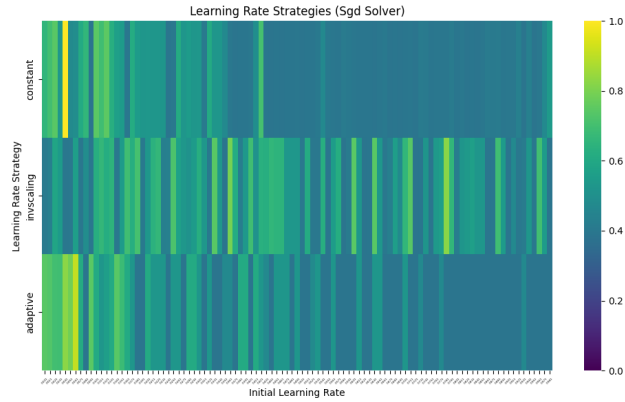


Figure 8: SGD performs best with low constant/adaptive

### 3.2.2 Effects of Solver Choice: ADAM



Figure 9: ADAM performs best with low invscaling/adaptive

## 4 Discussion

### 4.1 Layers and Units

#### 4.1.1 Units Per Layer

As seen in section 3.1.3, increasing units per layer negatively affects the neural network score unless the number of layers is also increased, and only up to a certain point. From Figure 2, you can see that any number of units per layer over

forty drastically affects the network score, regardless of the number of layers. Additionally, notice that at least four units are needed in the heat maps with higher classes for the network to perform well. This can be because more units are required in the output layer, so it would intuitively insinuate that a network with a shape similar to "i" is required. The network will perform better if the number of units in the hidden layers is closer to the number of units in the output layer. Thus, a network with a shape similar to "i" is not performant because, at the final layers, there are too few units feeding values into the output layer, spreading thin the information accumulated through the network.

## 4.2 Number of Layers

Looking at the number of layers, we notice a similar pattern: increasing numbers of layers with a good choice of units per layer does not have any meaningful effects (in the case of units, an ever-increasing number has adverse effects). However, notice that as the number of classes increases, a minimum number of layers is required for the network to perform decently, similar to the number of units. Putting this together, after a good choice of units per layer, the number of layers has diminishing returns. This can be explained by imagining that, given a good number of units per layer, adding more and more layers, will only generate weights with very slightly better values between each layer. Although the accuracy may improve, it will only be by a little.

## 4.3 Heatmap Jaggedness

To explain the jaggedness of the heatmaps, consider that not all networks had their optimizers converge. When training the networks, a maximum of 200 iterations was used. This means that certain combinations of the number of layers and units per layer may have required more iterations to have the optimizer converge or reach a good enough optimization. On the other hand, some networks may have reached an optimal value in the few iterations by chance. This explains how some networks with a few layers and many units per layer (predicted to be a bad network) could have performed well and vice-versa.

## 4.4 Solver Choice and Learning Rates

Based on the results of the tests conducted using various learning rates, learning rate algorithms, and solvers, various conclusions can be supported and improved with further testing. The results provide evidence that the choice regarding these parameters has a direct and strong effect on the effectiveness of the neural network.

### 4.4.1 Initial Learning Rate

Each test was run using 100 different learning rates spanning from 0.001 and 0.991. Generally speaking, under the specified parameters, it seems clear that



a lower initial learning rate provides more accuracy to the model than a higher one across both solvers and all learning rate strategies. A higher learning rate allows for faster convergence, but also creates a higher potential for overfitting the data. A lower learning rate creates a slower descent towards the correct prediction, which these results support to be a more effective approach. The only exception to this trend came from the SGD solver using the invscaling learning strategy; instead of having a higher accuracy with lower initial learning rates, the accuracy was scattered randomly across all values.

#### 4.4.2 Learning Rate Strategies: Constant

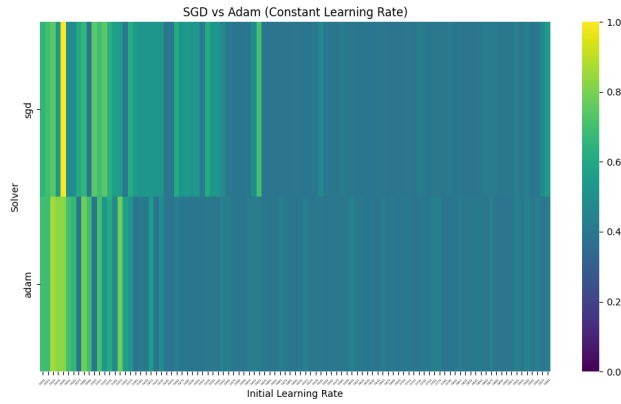


Figure 10: SGD generally performs better with these conditions

The constant learning rate uses a fixed learning rate for the duration of testing. For both solvers, there was a clear drop in accuracy as the initial learning rate increased. This intuitively makes sense, as an immutable high learning rate creates a big risk for overfitting. Although ADAM performed slightly more consistently across the lower learning rates, the SGD solver had a higher maximum accuracy rating.

#### 4.4.3 Learning Rate Strategies: Invscaling

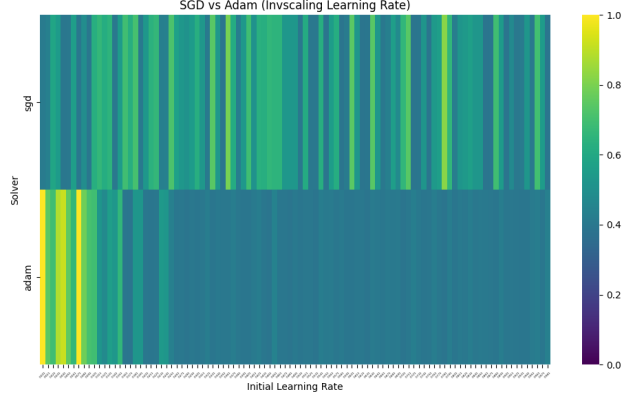


Figure 11: Adam has higher predictability

Invscaling uses a descending learning rate to allow the model to both move quickly towards convergence and reduce the risk of overfitting. The invscaling learning rate is defined as:

$$\text{learning rate}_t = \frac{\text{initial learning rate}}{(1 + \text{power\_t} \times t)^{\text{exponent}}} \quad (1)$$

where:

- **initial learning rate** is the starting learning rate.
- **t** is the current iteration or epoch number.
- **power\_t** is a constant parameter that scales the rate of decrease.
- **exponent** controls the decay rate, often set to 0.5.

Using this equation, the ADAM solver proved to be more efficient than the results shown using the constant learning rate; it was significantly more effective at lower learning rates, but was still equally inconsistent with higher learning rates. However, the SGD solver created a very randomized spread across all learning rates. This is an interesting observation and demonstrates the power invscaling can have to combat overfitting data.

#### 4.4.4 Learning Rate Strategies: Adaptive



Figure 12: Adam has higher consistency across some rates

The adaptive learning rate strategy allows for a more customized approach to the changing learning rate using these properties:

- The **Learning Rate Adjustment** changes based on the magnitudes calculated during training. This allows the learning rate to be both increased or decreased based on the previous step's success rate.
- **Gradient Accumulation:** The adaptive learning rate keeps track of past gradients to help scale the learning rates based on how frequently that node has been updated.
- **Per-Parameter Learning Rate:** The adaptive learning rate uses a separate learning rate for each parameter in the model, allowing for maximum flexibility and potential for success.

Using this strategy, it seems there is a better level of consistency across the board for all learning rates. Although ADAM performed slightly better under certain learning rates using the invscaling method, there is a slight improvement across higher learning rates.

## 5 Conclusion

Based on the conclusions drawn, the process of choosing an optimal number of layers and units per layer for a network and a given dataset can be outlined. One can start with at least as many units as classes, and progressively increase the number of layers until a balance between improved accuracy and training time is achieved. Although the training time for each network was tested, the data was

not included due to the difficulty of representing the four dimensions. However, this information would be valuable in making a more informed decision.

Overall, it seems clear a lower initial learning rate proves to be more consistent than a higher value across both solvers and all descent strategies. The ADAM solver generally has higher maximum accuracy than the SGD solver, but its accuracy drops off much faster with increasing rates. Using these parameters, the data suggests the most effective combination of solvers and learning rate strategies is the ADAM solver with an initial learning rate below 0.1 and either the invscaling or adaptive strategy. Intuitively, the adaptive strategy should work more consistently across randomized data and testing due to its greater flexibility.

## 6 Code

[GitHub Repository](#)

### 6.1 ryan.py

```
1
2 import warnings
3 from time import time
4
5 import matplotlib.pyplot as plt
6 import numpy as np
7 import seaborn as sns
8 from sklearn.exceptions import ConvergenceWarning
9 from sklearn.model_selection import train_test_split
10 from sklearn.neural_network import MLPClassifier
11 from tqdm import tqdm
12
13 from data_generators import aniso, blob
14 from read_and_write import write_to_file
15
16 # Ignoring convergence warnings to not mess up the progress bar
17 warnings.filterwarnings("ignore", category=ConvergenceWarning)
18
19 # Default values
20 DATA_SIZE = 1000
21 NUMBER_OF_FEATURES = 2
22 RANDOM_STATE = 1
23
24 # Dataset ranges
25 DATA_TYPES = [blob, aniso]
26 MIN_NUMBER_OF_CLASSES = 2
27 MAX_NUMBER_OF_CLASSES = 3
28
29 # Network architecture fixed values
```

```

30 NUMBER_OF_LAYERS = 10      # 10+ layers based on results of network
    architecture tests
31 NUMBER_OF_UNITS = 5       # Lower units(1-5) based on results of network
    architecture tests
32
33 # Learning rate ranges and solvers
34 MIN_LEARNING_RATE_INIT = 0.001
35 MAX_LEARNING_RATE_INIT = 0.991
36 LEARNING_RATE_STEP_SIZE = 0.01
37 SOLVERS = ['sgd', 'adam']
38 LEARNING_RATE_STRATEGIES = ['constant', 'invscaling', 'adaptive']
39
40 # Operations counts
41 total_datasets = len(DATA_TYPES) * (
42     MAX_NUMBER_OF_CLASSES - MIN_NUMBER_OF_CLASSES + 1
43 )
44 total_networks = len(SOLVERS) * len(LEARNING_RATE_STRATEGIES) * (
45     int((MAX_LEARNING_RATE_INIT - MIN_LEARNING_RATE_INIT) /
46         LEARNING_RATE_STEP_SIZE) + 1
47 )
48 # Initializing timer
49 master_start_time = time()
50
51 # Building datasets
52 data_pbar = tqdm(total=total_datasets, desc="Building Datasets")
53
54 datasets = []
55 for data_type in DATA_TYPES:
56     for number_of_classes in range(MIN_NUMBER_OF_CLASSES,
57                                     MAX_NUMBER_OF_CLASSES + 1):
58         X, y = data_type(DATA_SIZE, number_of_classes,
59                           NUMBER_OF_FEATURES, RANDOM_STATE)
60         X_train, X_test, y_train, y_test = train_test_split(X, y)
61         datasets.append((data_type.__name__, number_of_classes, X_train,
62                           X_test, y_train, y_test))
63         data_pbar.update(1)
64
65 data_pbar.close()
66
67 # Running experiments
68 experiment_pbar = tqdm(total=total_networks * total_datasets,
69                        desc="Running Experiments")
70
71 results = []
72 for solver in SOLVERS:
73     for learning_rate_strategy in LEARNING_RATE_STRATEGIES:
74         for learning_rate_init in np.arange(MIN_LEARNING_RATE_INIT,
75                                             MAX_LEARNING_RATE_INIT,
76                                             LEARNING_RATE_STEP_SIZE):
77             network = MLPClassifier(

```

```

72         hidden_layer_sizes=(NUMBER_OF_UNITS,) * NUMBER_OF_LAYERS,
73         solver=solver,
74         learning_rate=learning_rate_strategy,
75         learning_rate_init=learning_rate_init
76     )
77
78     for dataset in datasets:
79         start_time = time()
80
81         X_train, X_test, y_train, y_test = dataset[2],
dataset[3], dataset[4], dataset[5]
82         network.fit(X_train, y_train)
83         score = network.score(X_test, y_test)
84
85         results.append((
86             dataset[0], dataset[1], solver,
learning_rate_strategy,
87             learning_rate_init, score, time() - start_time
88         ))
89
90         experiment_pbar.update(1)
91
92 experiment_pbar.close()
93 print(f"Total time: {time() - master_start_time}")
94
95 # Writing results to file in case of somehow losing the data
96 write_to_file(results, "solver_learning_rate_results.txt")
97
98 # Converting results to a numpy array for easier processing
99 results = np.array(results, dtype=[
100     ("type_of_data", "U50"),
101     ("classes", int),
102     ("solver", "U10"),
103     ("learning_rate_strategy", "U10"),
104     ("learning_rate_init", float),
105     ("score", float),
106     ("total_time", float),
107 ])
108
109 # Aggregate results for blobs and aniso
110 def aggregate_results(results):
111     aggregated_results = []
112     for solver in SOLVERS:
113         for lr_strategy in LEARNING_RATE_STRATEGIES:
114             for learning_rate_init in np.arange(MIN_LEARNING_RATE_INIT,
MAX_LEARNING_RATE_INIT, LEARNING_RATE_STEP_SIZE):
115                 mask = (results["solver"] == solver) &
(results["learning_rate_strategy"] == lr_strategy)
116                 scores_for_lr_init = results[mask &
(results["learning_rate_init"] == learning_rate_init)]["score"]

```

```

117         mean_score = np.mean(scores_for_lr_init)
118         aggregated_results.append((solver, lr_strategy,
119         learning_rate_init, mean_score))
119     return np.array(aggregated_results, dtype=[("solver", "U10"),
120         ("learning_rate_strategy", "U10"), ("learning_rate_init", float),
121         ("score", float)])
120
121 aggregated_results = aggregate_results(results)
122
123 # Heatmap function
124 def plot_heatmap(data, row_labels, col_labels, title, xlabel, ylabel):
125     plt.figure(figsize=(10, 6))
126     sns.heatmap(data, xticklabels=col_labels, yticklabels=row_labels,
127         cmap="viridis", cbar=True, vmin=0, vmax=1)
128     plt.title(title)
129     plt.xlabel(xlabel)
130     plt.ylabel(ylabel)
131     plt.xticks(rotation=45, fontsize=3)
132     plt.tight_layout()
133     plt.show()
134
135 # Heatmap for sgd vs adam for each learning rate strategy
136 def plot_solver_comparison(aggregated_results, lr_strategy):
137     heatmap_data = np.zeros((2, len(np.arange(MIN_LEARNING_RATE_INIT,
138         MAX_LEARNING_RATE_INIT, LEARNING_RATE_STEP_SIZE))))
139     solvers = ['sgd', 'adam']
140
141     for i, solver in enumerate(solvers):
142         for j, learning_rate_init in
143             enumerate(np.arange(MIN_LEARNING_RATE_INIT, MAX_LEARNING_RATE_INIT,
144                 LEARNING_RATE_STEP_SIZE)):
145             mask = (aggregated_results["solver"] == solver) &
146                 (aggregated_results["learning_rate_strategy"] == lr_strategy)
147             score = aggregated_results[mask &
148                 (aggregated_results["learning_rate_init"] ==
149                 learning_rate_init)]["score"]
150             heatmap_data[i, j] = np.mean(score)
151
152     plot_heatmap(heatmap_data, solvers,
153         np.round(np.arange(MIN_LEARNING_RATE_INIT, MAX_LEARNING_RATE_INIT,
154             LEARNING_RATE_STEP_SIZE), 3),
155         f"SGD vs Adam ({lr_strategy.capitalize()} Learning
156         Rate)", "Initial Learning Rate", "Solver")
157
158 # Heatmap for comparing learning rate strategies for each solver
159 def plot_lr_strategy_comparison(aggregated_results, solver):
160     heatmap_data = np.zeros((3, len(np.arange(MIN_LEARNING_RATE_INIT,
161         MAX_LEARNING_RATE_INIT, LEARNING_RATE_STEP_SIZE))))
162     lr_strategies = ['constant', 'invscaling', 'adaptive']

```

```

153     for i, lr_strategy in enumerate(lr_strategies):
154         for j, learning_rate_init in
            enumerate(np.arange(MIN_LEARNING_RATE_INIT, MAX_LEARNING_RATE_INIT,
            LEARNING_RATE_STEP_SIZE)):
155             mask = (aggregated_results["solver"] == solver) &
            (aggregated_results["learning_rate_strategy"] == lr_strategy)
156             score = aggregated_results[mask &
            (aggregated_results["learning_rate_init"] ==
            learning_rate_init)]["score"]
157             heatmap_data[i, j] = np.mean(score)
158
159     plot_heatmap(heatmap_data, lr_strategies,
            np.round(np.arange(MIN_LEARNING_RATE_INIT, MAX_LEARNING_RATE_INIT,
            LEARNING_RATE_STEP_SIZE), 3),
160                 f"Learning Rate Strategies ({solver.capitalize()}
            Solver)", "Initial Learning Rate", "Learning Rate Strategy")
161
162     # Plot visualizations
163     plot_solver_comparison(aggregated_results, 'constant')
164     plot_solver_comparison(aggregated_results, 'invscaling')
165     plot_solver_comparison(aggregated_results, 'adaptive')
166
167     plot_lr_strategy_comparison(aggregated_results, 'sgd')
168     plot_lr_strategy_comparison(aggregated_results, 'adam')

```

## 6.2 rafael.py

```

1     """
2     Executes experiments to analyze the effects of different number of
            layers and units per layer on the score of a network.
3     """
4
5     from os import makedirs, path
6     from time import time
7     from warnings import filterwarnings
8
9     import matplotlib.pyplot as plt
10    from numpy import array as nparray
11    from numpy import zeros as npzeros
12    from sklearn.exceptions import ConvergenceWarning
13    from sklearn.model_selection import train_test_split
14    from sklearn.neural_network import MLPClassifier
15    from tqdm import tqdm
16
17    from data_generators import aniso, blob
18    from read_and_write import write_to_file
19
20    # ignoring convergence warnings to not mess up the progress bar

```



```

21 filterwarnings("ignore", category=ConvergenceWarning)
22
23 # default values
24 FOLDER_NAME = "layers_units"
25 DATA_SIZE = 500 # barely any difference between 500-5000 data points
26 NUMBER_OF_FEATURES = 2
27 RANDOM_STATE = 1
28
29 print("Press enter for default values")
30
31 # dataset ranges
32 DATA_TYPES = [blob, aniso]
33 MIN_NUMBER_OF_CLASSES = int(input("MIN_NUMBER_OF_CLASSES or 2: ") or 2)
34 MAX_NUMBER_OF_CLASSES = int(input("MAX_NUMBER_OF_CLASSES or 7: ") or 7)
35
36 # network architecture ranges
37 MIN_NUMBER_OF_LAYERS = int(input("MIN_NUMBER_OF_LAYERS or 1: ") or 1)
38 MAX_NUMBER_OF_LAYERS = int(input("MAX_NUMBER_OF_LAYERS or 40: ") or 40)
39 MIN_NUMBER_OF_UNITS = int(input("MIN_NUMBER_OF_UNITS or 1: ") or 1)
40 MAX_NUMBER_OF_UNITS = int(input("MAX_NUMBER_OF_UNITS or 40: ") or 40)
41
42 # operations counts
43 TOTAL_DATASETS = len(DATA_TYPES) * (
44     MAX_NUMBER_OF_CLASSES - MIN_NUMBER_OF_CLASSES + 1
45 )
46 TOTAL_NETWORKS = (MAX_NUMBER_OF_LAYERS - MIN_NUMBER_OF_LAYERS + 1) * (
47     MAX_NUMBER_OF_UNITS - MIN_NUMBER_OF_UNITS + 1
48 )
49
50
51 def generate_datasets() -> list[dict]:
52     pbar = tqdm(total=TOTAL_DATASETS, desc="Generate Datasets")
53
54     datasets = []
55     for data_type in DATA_TYPES:
56         for number_of_classes in range(
57             MIN_NUMBER_OF_CLASSES, MAX_NUMBER_OF_CLASSES + 1
58         ):
59             X, y = data_type(
60                 DATA_SIZE,
61                 number_of_classes,
62                 NUMBER_OF_FEATURES,
63                 RANDOM_STATE,
64             )
65
66             X_train, X_test, y_train, y_test = train_test_split(X, y)
67
68             datasets.append(
69                 {
70                     "data_type": data_type.__name__,

```

```

71         "number_of_classes": number_of_classes,
72         "X_train": X_train,
73         "X_test": X_test,
74         "y_train": y_train,
75         "y_test": y_test,
76     }
77 )
78 pbar.update(1)
79
80 pbar.close()
81
82 return datasets
83
84
85 def run_experiments(datasets: list[dict]) -> list[tuple]:
86     pbar = tqdm(total=TOTAL_NETWORKS * TOTAL_DATASETS, desc="Run
87         Experiments")
88
89     results = []
90     for layers in range(MIN_NUMBER_OF_LAYERS, MAX_NUMBER_OF_LAYERS + 1):
91         for units in range(MIN_NUMBER_OF_UNITS, MAX_NUMBER_OF_UNITS + 1):
92             network = MLPClassifier(hidden_layer_sizes=(units,) * layers)
93
94             for dataset in datasets:
95                 start_time = time()
96
97                 X_train, X_test, y_train, y_test = (
98                     dataset["X_train"],
99                     dataset["X_test"],
100                     dataset["y_train"],
101                     dataset["y_test"],
102                 )
103
104                 # train network and get score
105                 network.fit(X_train, y_train)
106                 score = network.score(X_test, y_test)
107
108                 results.append(
109                     (
110                         dataset["data_type"],
111                         dataset["number_of_classes"],
112                         layers,
113                         units,
114                         score,
115                         time() - start_time,
116                     )
117                 )
118
119                 pbar.update(1)

```

```

120     pbar.close()
121
122     return results
123
124
125 def download_heat_maps(results: list[tuple]) -> None:
126     # converting results to nparray to easily access columns
127     results = nparray(
128         results,
129         dtype=[
130             ("type_of_data", "U50"),
131             ("classes", int),
132             ("layers", int),
133             ("units", int),
134             ("score", float),
135             ("total_time", float),
136         ],
137     )
138
139     for data_type in DATA_TYPES:
140         for number_of_classes in range(
141             MIN_NUMBER_OF_CLASSES, MAX_NUMBER_OF_CLASSES + 1
142         ):
143             data_type_name = data_type.__name__
144             # processing data
145             mask = (results["type_of_data"] == data_type_name) & (
146                 results["classes"] == number_of_classes
147             )
148             filtered_results = results[mask]
149
150             layers = filtered_results["layers"]
151             units = filtered_results["units"]
152             score = filtered_results["score"]
153
154             score_2d = npzeros(
155                 (
156                     MAX_NUMBER_OF_LAYERS - MIN_NUMBER_OF_LAYERS + 1,
157                     MAX_NUMBER_OF_UNITS - MIN_NUMBER_OF_UNITS + 1,
158                 )
159             )
160
161             for l, u, s in zip(layers, units, score):
162                 score_2d[l - MIN_NUMBER_OF_LAYERS, u -
MIN_NUMBER_OF_UNITS] = s
163
164             # setting up plot
165             aspect_ratio = score_2d.shape[1] / score_2d.shape[0]
166             fig_width = 10
167             fig_height = fig_width / aspect_ratio
168

```

```

169 plt.figure(figsize=(fig_width, fig_height), dpi=300)
170
171 im = plt.imshow(
172     score_2d,
173     cmap="magma",
174     extent=(
175         MIN_NUMBER_OF_LAYERS - 0.5,
176         MAX_NUMBER_OF_LAYERS + 0.5,
177         MIN_NUMBER_OF_UNITS - 0.5,
178         MAX_NUMBER_OF_UNITS + 0.5,
179     ),
180     origin="lower",
181     aspect="auto",
182     vmin=0,
183     vmax=1,
184 )
185
186 cbar = plt.colorbar(im, label="Score")
187 plt.xlabel("Number of Layers")
188 plt.ylabel("Number of Units")
189 plt.title(
190     f"Score vs. Number of Layers and Units
191     ({data_type_name}, {number_of_classes} Classes)"
192 )
193
194 plt.xticks(range(MIN_NUMBER_OF_LAYERS, MAX_NUMBER_OF_LAYERS
195 + 1))
196 plt.yticks(range(MIN_NUMBER_OF_UNITS, MAX_NUMBER_OF_UNITS +
197 1))
198
199 plt.tight_layout()
200
201 # Downloading plots
202 makedirs(FOLDER_NAME, exist_ok=True)
203
204 filename =
205 f"size{DATA_SIZE}_features{NUMBER_OF_FEATURES}_random{RANDOM_STATE}
206 _type{data_type_name}_classes{number_of_classes}.png"
207 filepath = path.join(FOLDER_NAME, filename)
208
209 plt.savefig(filepath, dpi=300, bbox_inches="tight")
210 plt.clf()
211 plt.close()
212
213 print(f"Heatmap saved as {filepath}")
214
215 # main
216 master_start_time = time()

```

```

215
216 datasets = generate_datasets()
217 results = run_experiments(datasets)
218
219 print(f"Total time: {time() - master_start_time}")
220
221 # saving data in any case
222 write_to_file(results, "rafael_results.txt")
223
224 download_heat_maps(results)

```

### 6.3 data\_generators.py

```

1  from typing import Tuple
2
3  import numpy as np
4  from sklearn.datasets import make_blobs
5
6
7  def blob(
8      samples: int, centers: int, features: int, random_state: int
9  ) -> Tuple[np.ndarray, np.ndarray]:
10     """Generates a blob dataset with labels"""
11
12     X, Y = make_blobs(
13         n_samples=samples,
14         centers=centers,
15         n_features=features,
16         random_state=random_state,
17     )
18     return X, Y
19
20
21 def aniso(
22     samples: int, centers: int, features: int, random_state: int
23 ) -> Tuple[np.ndarray, np.ndarray]:
24     """Generates a aniso dataset with labels"""
25
26     X, Y = blob(samples, centers, features, random_state)
27     transformation = [[0.6, -0.6], [-0.4, 0.8]]
28     X = np.dot(X, transformation)
29     return X, Y

```

### 6.4 read\_and\_write.py

```

1  from typing import Any, List
2

```

```

3
4 def write_to_file(arr: List[List[Any]], file_name: str) -> None:
5     """Writes a 2D array to file_name with columns separated with tabs
6     and rows with newlines"""
7
8     with open(file_name, "w") as f:
9         for row in arr:
10             f.write("\t".join(map(str, row)) + "\n")
11
12 def read_to_array(file_name: str) -> List[List[Any]]:
13     """Reads a file and returns a 2D array"""
14
15     with open(file_name, "r") as f:
16         return [line.strip().split("\t") for line in f]

```